

# LLM supporting tools that make the ecosystem

*Paul Rodriguez*  
*San Diego Supercomputer Center*

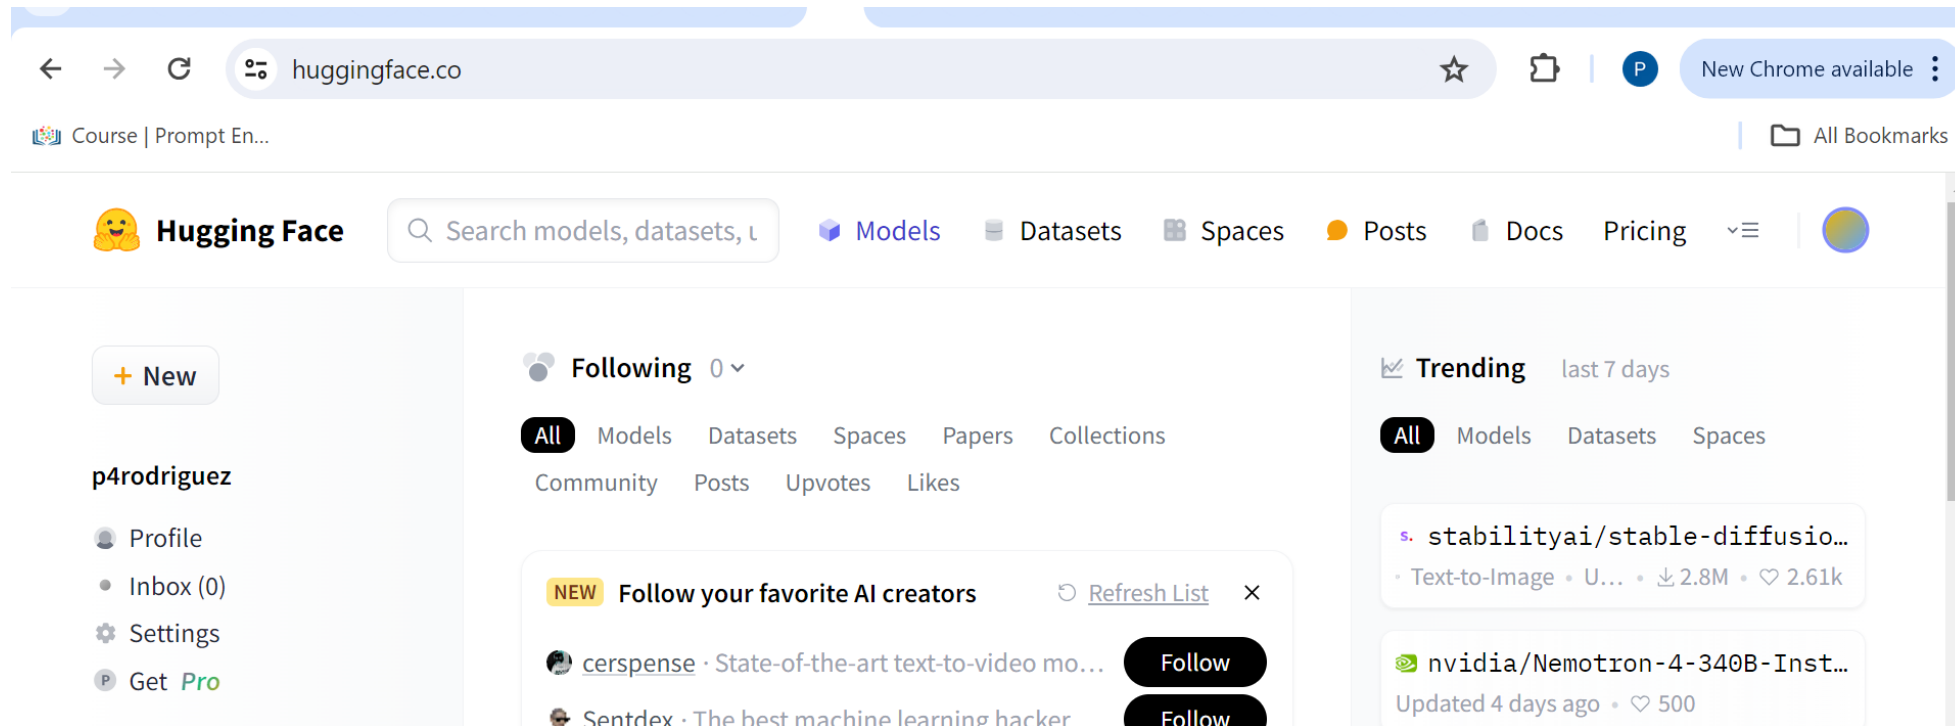
*2024 CIML Summer Institute*

# Outline

- **The Hugging Face ecosystem**
- **Langchain/Langgraph**
- **Using Hugging Face tutorials discussion/demo**

# Hugging Face Hub

- **huggingface.com** is a hub of models, data, tutorials for using AI models



# Using Hugging Face (HF)

- HF provides python packages to make it (relatively) easy to run models
- Accessing models and/or data requires an HF authentication token
- Some of the main packages are:
  - pipeline: to run inferencing
  - diffusers: for diffusion models
  - transformers: for LLMs
  - accelerate: for efficient and/or parallel execution
  - datasets: to access data from hub

# Hugging Face Abstractions

- HF package are intended to be more abstract than a Google, Meta, Msoft, etc. package, and work the same for all of them
- For example, the pipeline function is built on more basic functions:

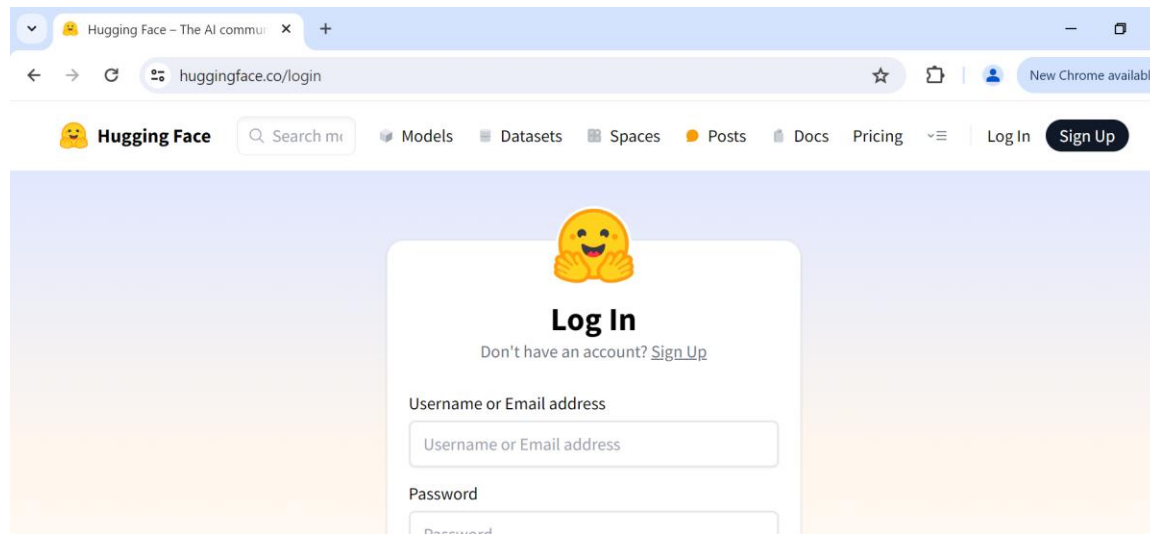
get “config”            (to get the model architecture)  
get “model”            (to download a pretrained model)  
get “tokenizer”        (to get the appropriate tokenizer)

- These can be invoked more directly if you want to do a pre-training or fine-tuning implementation, or experimentation

# Getting account set up

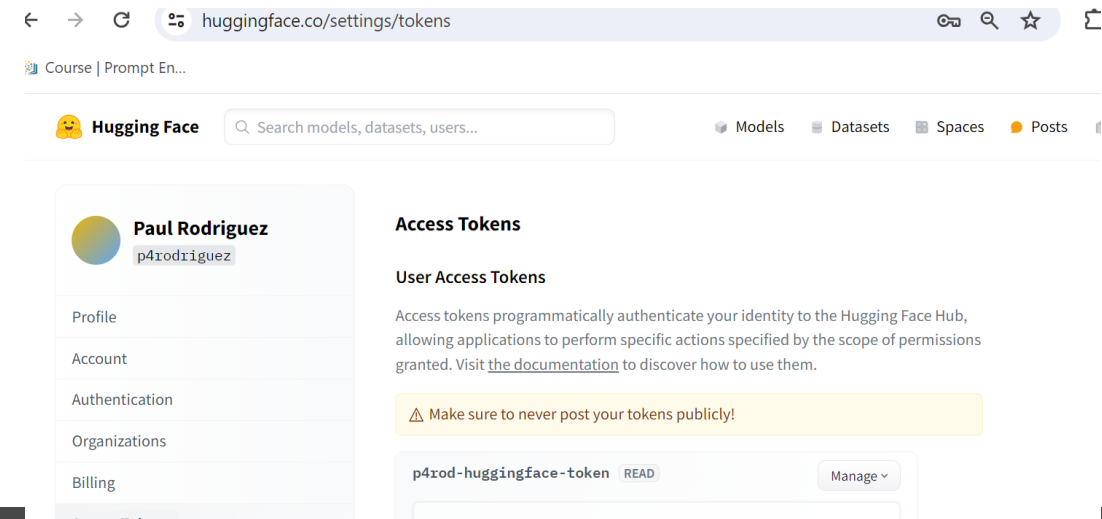
- Create an account on huggingface and get authentication token

huggingface.co/login



The screenshot shows the Hugging Face login page in a web browser. The browser's address bar displays 'huggingface.co/login'. The page features a 'Log In' form with a yellow smiley face icon at the top. Below the icon, the text 'Log In' is displayed, followed by a link 'Don't have an account? Sign Up'. The form has two input fields: 'Username or Email address' and 'Password'. The navigation bar at the top includes the Hugging Face logo, a search bar, and links for Models, Datasets, Spaces, Posts, Docs, Pricing, Log In, and Sign Up.

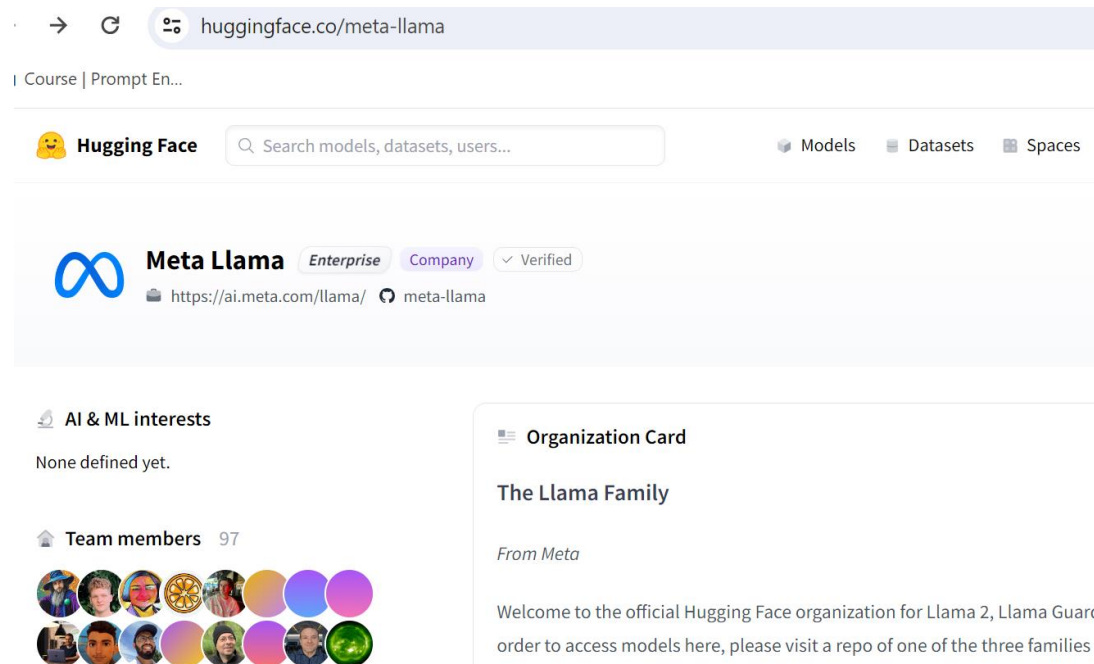
Account -> settings -> get token



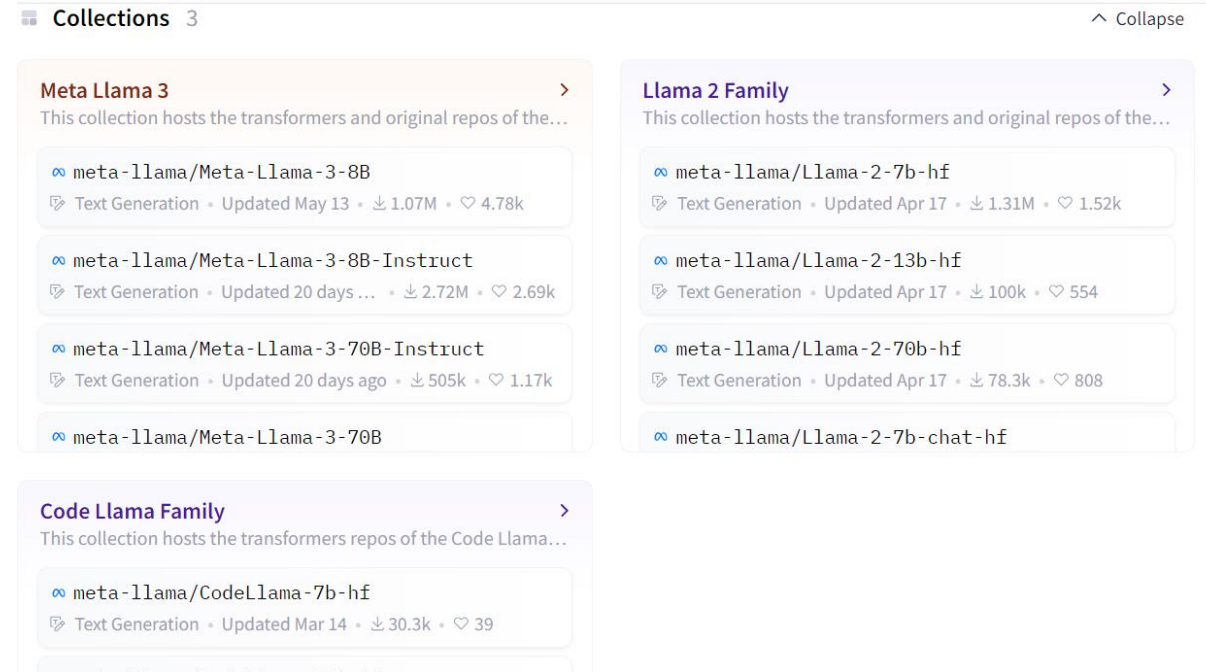
The screenshot shows the Hugging Face settings page, specifically the 'Access Tokens' section. The browser's address bar displays 'huggingface.co/settings/tokens'. The page features a sidebar with the user's profile 'Paul Rodriguez' and a list of settings: Profile, Account, Authentication, Organizations, and Billing. The main content area is titled 'Access Tokens' and includes a section for 'User Access Tokens'. It explains that access tokens programmatically authenticate identity to the Hugging Face Hub. A warning message states: 'Make sure to never post your tokens publicly!'. Below this, a table lists the user's access tokens, showing one token named 'p4xrod-huggingface-token' with a 'READ' scope and a 'Manage' button.

# Hugging Face model repo

- **Example:** Meta has a family of Llama models that vary by size, response training, 'hf' format, release date, etc..



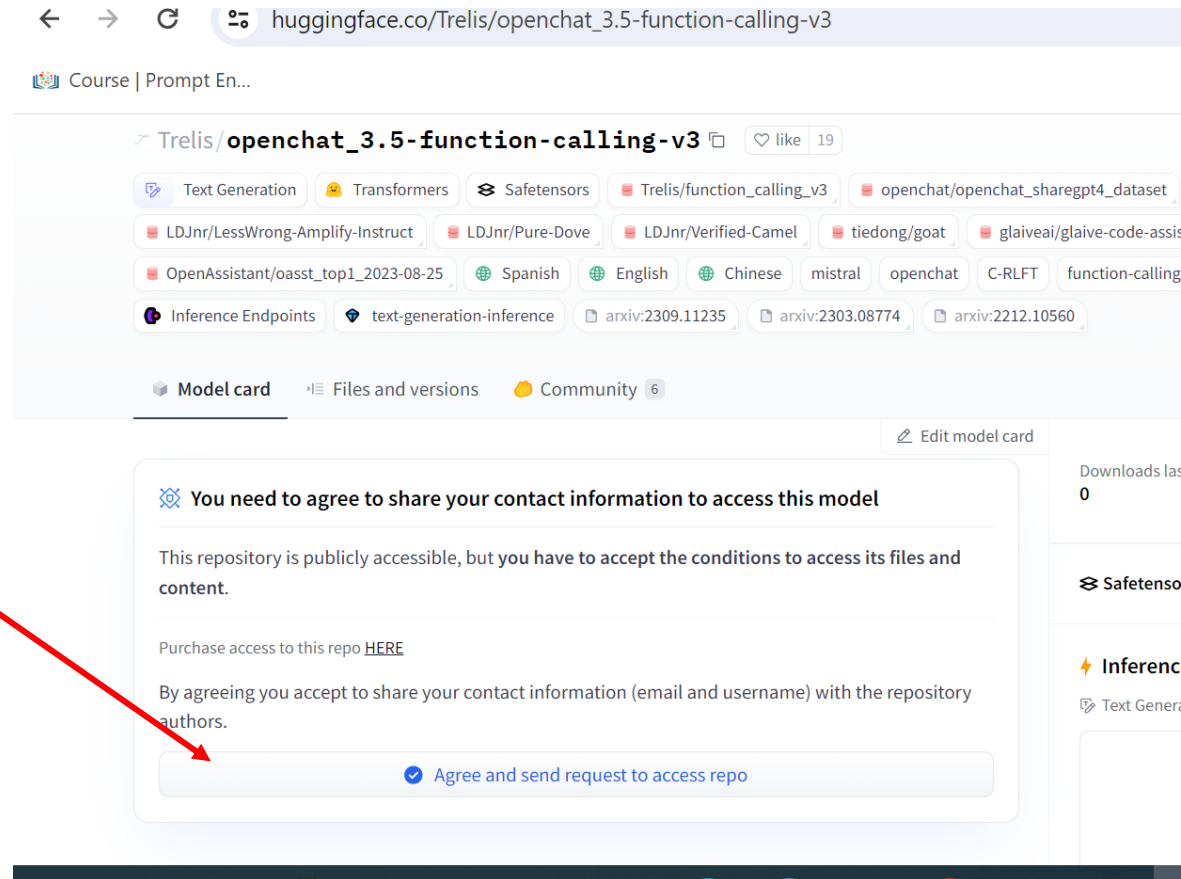
The screenshot shows the Hugging Face profile page for Meta Llama. The browser address bar displays 'huggingface.co/meta-llama'. The page header includes the Hugging Face logo, a search bar, and navigation links for Models, Datasets, and Spaces. The main profile section features the Meta Llama logo, the name 'Meta Llama', and tags for 'Enterprise', 'Company', and 'Verified'. Below this, there are links to the official website and the GitHub repository. The page also includes sections for 'AI & ML interests' (currently none defined) and 'Team members' (97 members, shown as a grid of profile pictures). An 'Organization Card' is visible, titled 'The Llama Family', with a welcome message for users.



The screenshot shows the 'Collections' page for Meta Llama on Hugging Face. The page is titled 'Collections 3' and has a 'Collapse' button. It lists three collections: 'Meta Llama 3', 'Llama 2 Family', and 'Code Llama Family'. Each collection is described as hosting transformers and original repos. The 'Meta Llama 3' collection lists four models: 'meta-llama/Meta-Llama-3-8B', 'meta-llama/Meta-Llama-3-8B-Instruct', 'meta-llama/Meta-Llama-3-70B-Instruct', and 'meta-llama/Meta-Llama-3-70B'. The 'Llama 2 Family' collection lists four models: 'meta-llama/Llama-2-7b-hf', 'meta-llama/Llama-2-13b-hf', 'meta-llama/Llama-2-70b-hf', and 'meta-llama/Llama-2-7b-chat-hf'. The 'Code Llama Family' collection lists one model: 'meta-llama/CodeLlama-7b-hf'. Each model entry includes a brief description, the number of downloads, and the number of likes.

# Hugging Face model repo

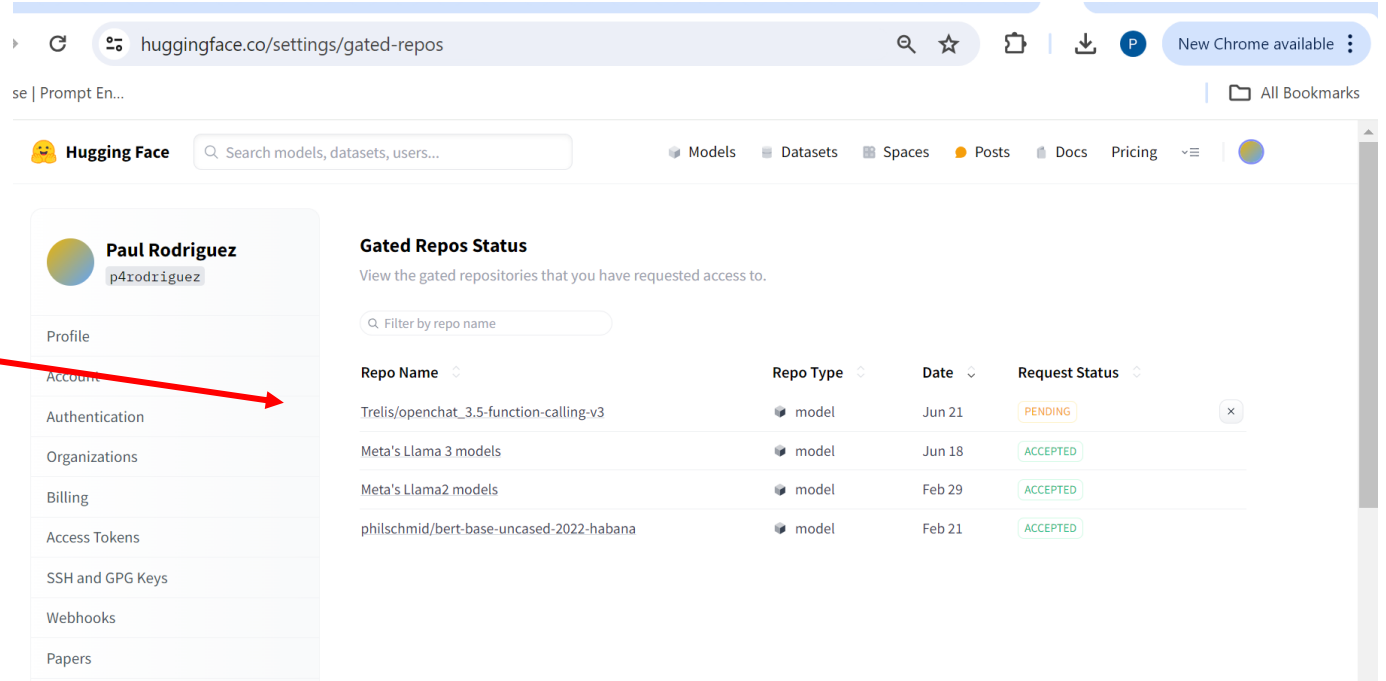
- Some models or data require that you request access here





# Hugging Face model repo

- Some models or data require that you request access
- Then go to account->setting->gated repo



The screenshot shows the Hugging Face website interface. The user is logged in as Paul Rodriguez (p4rodriguez). The left sidebar contains a menu with items: Profile, Account, Authentication, Organizations, Billing, Access Tokens, SSH and GPG Keys, Webhooks, and Papers. The 'Account' item is highlighted with a red arrow pointing to the 'Gated Repos Status' section on the right.

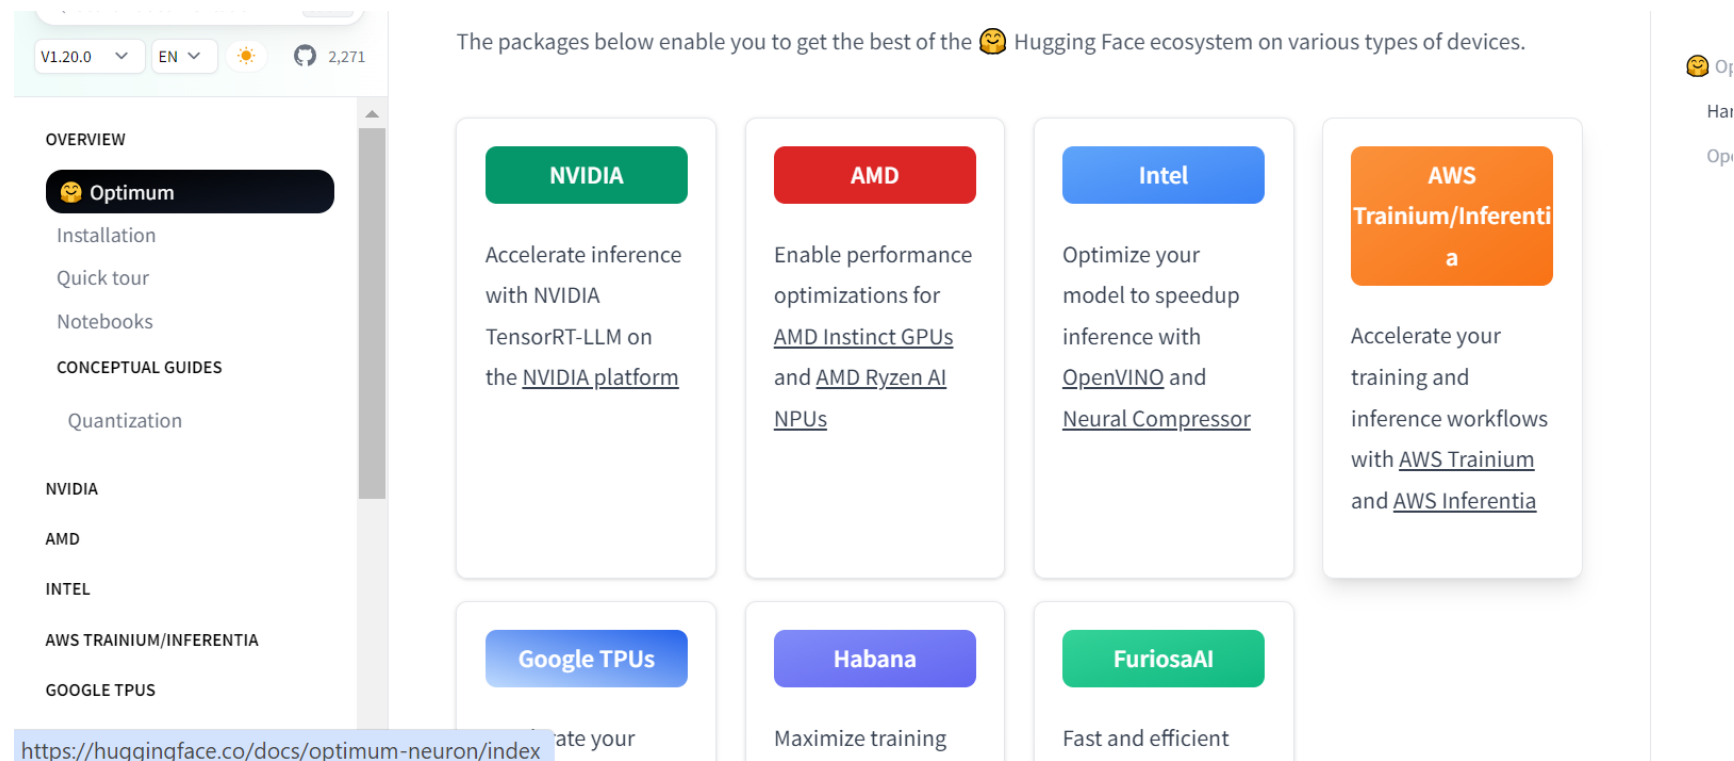
**Gated Repos Status**  
View the gated repositories that you have requested access to.

Filter by repo name

Repo Name	Repo Type	Date	Request Status
Trelis/openchat_3.5-function-calling-v3	model	Jun 21	PENDING
Meta's Llama 3 models	model	Jun 18	ACCEPTED
Meta's Llama2 models	model	Feb 29	ACCEPTED
philschmid/bert-base-uncased-2022-habana	model	Feb 21	ACCEPTED

# Hardware Partners

- Hugging Face also incorporates accelerator libraries into their packages



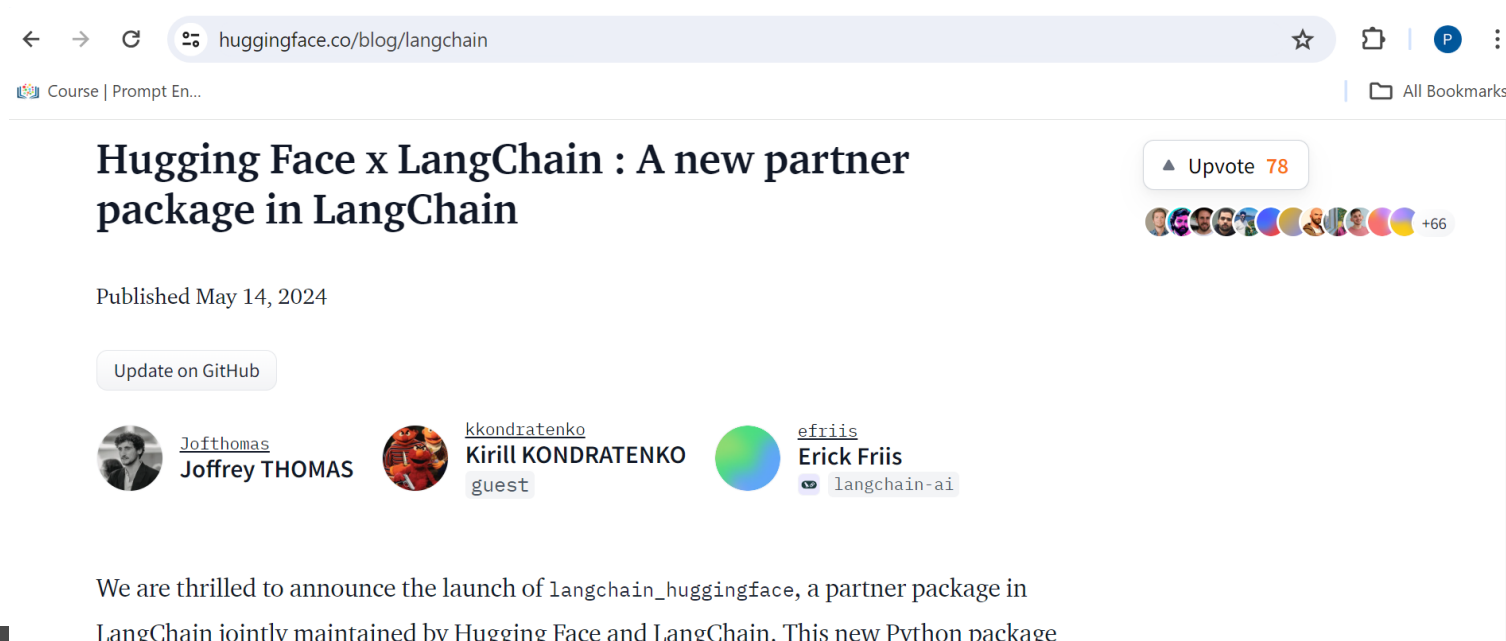
The screenshot displays the Hugging Face Optimum documentation page. The left sidebar contains a navigation menu with sections: OVERVIEW (with 'Optimum' selected), Installation, Quick tour, Notebooks, CONCEPTUAL GUIDES (with 'Quantization' selected), NVIDIA, AMD, INTEL, AWS TRAINIUM/INFERENCEIA, and GOOGLE TPUS. The main content area features a header stating, 'The packages below enable you to get the best of the 🤗 Hugging Face ecosystem on various types of devices.' Below this, there are eight cards representing different hardware partners:

- NVIDIA**: Accelerate inference with NVIDIA TensorRT-LLM on the [NVIDIA platform](#)
- AMD**: Enable performance optimizations for [AMD Instinct GPUs](#) and [AMD Ryzen AI NPU](#)s
- Intel**: Optimize your model to speedup inference with [OpenVINO](#) and [Neural Compressor](#)
- AWS Trainium/Inferentia**: Accelerate your training and inference workflows with [AWS Trainium](#) and [AWS Inferentia](#)
- Google TPUs**: Accelerate your training
- Habana**: Maximize training
- FuriosaAI**: Fast and efficient

A URL is visible at the bottom left: <https://huggingface.co/docs/optimum-neuron/index>

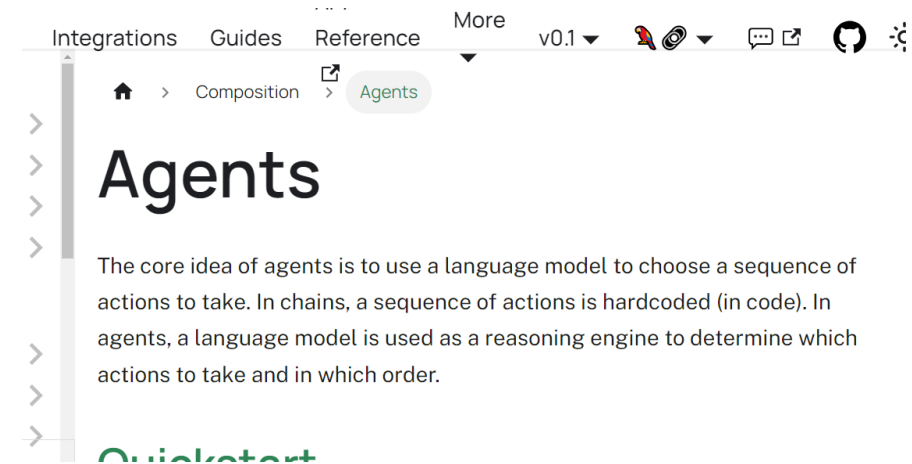
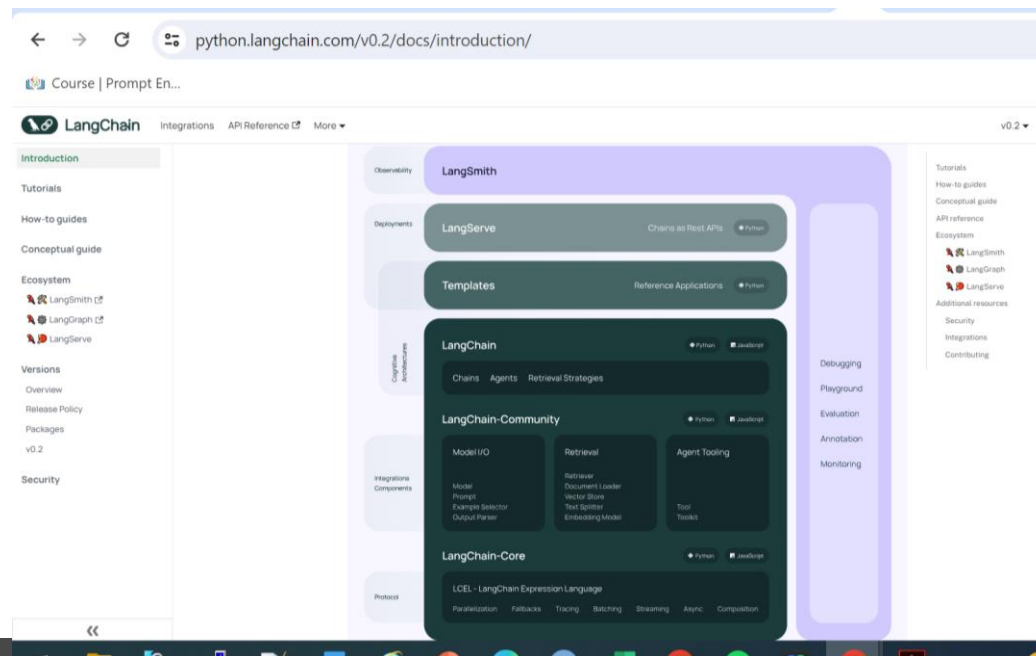
# Partner packages

- Hugging Face incorporates other packages
- For example, Langchain helps run RAG applications by:
  - accessing PDF files or URL text as raw documents
  - creating database of documents (split into chunks and vectorized)
  - Setting up prompts that include relevant contexts



# LangChain/LangGraph:

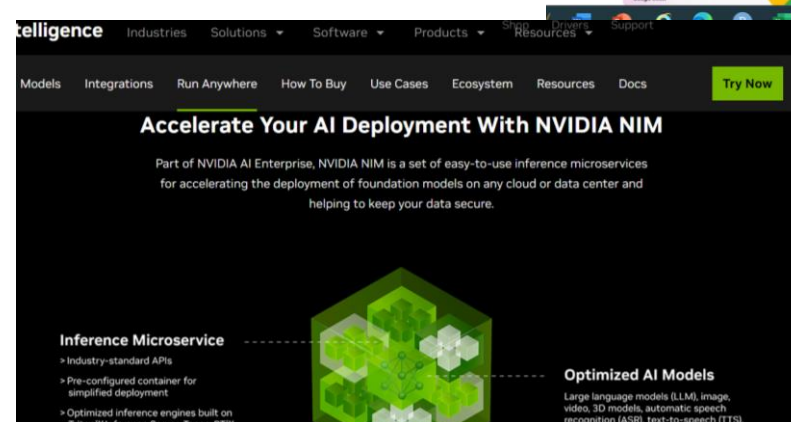
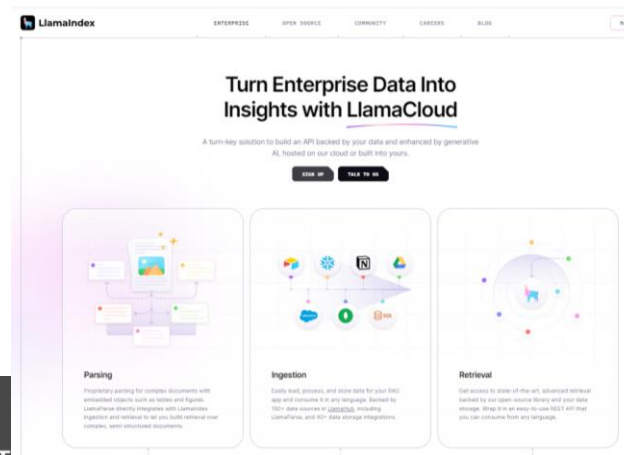
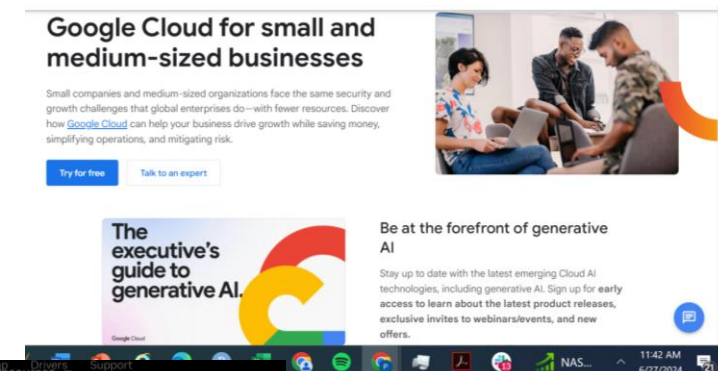
- Not an LLM provider (like openAI, or Hugging Face) but provides a standard interface
- Has ecosystem of tools beyond RAG – e.g. Agents: “a language model is used as a reasoning engine to determine which actions to take and in which order.”



# Other ecosystems to run LLMs

Many Big Tech firms have LLM services that provide ‘turn key’ solutions/tools for RAG, with APIs, open models, interfaces, deployment., etc.. geared for businesses.

- Nvidia (NeMo)
- Google (cloud) AI
- LlamaIndex (meta)
- Etc...



# Keras extension for NLP

- KerasNLP also has models, data, functions to run different LLMs (no tokens)
- Includes pre-trained LLMs base or full model, for example:
  - GPT2Backbone the model without task specific output layers
  - GPT2CausalLM the model with output predictions
  - GPT2CausalLMPreprocessor the preprocessor that feeds model.fit

# Expanse Demo/Exercise setting up HuggingFace and Langchain

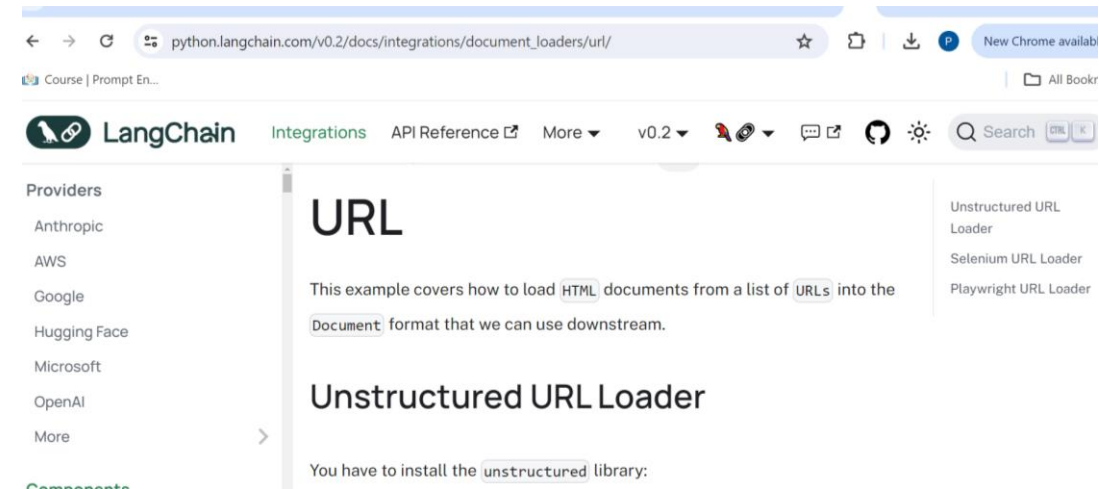
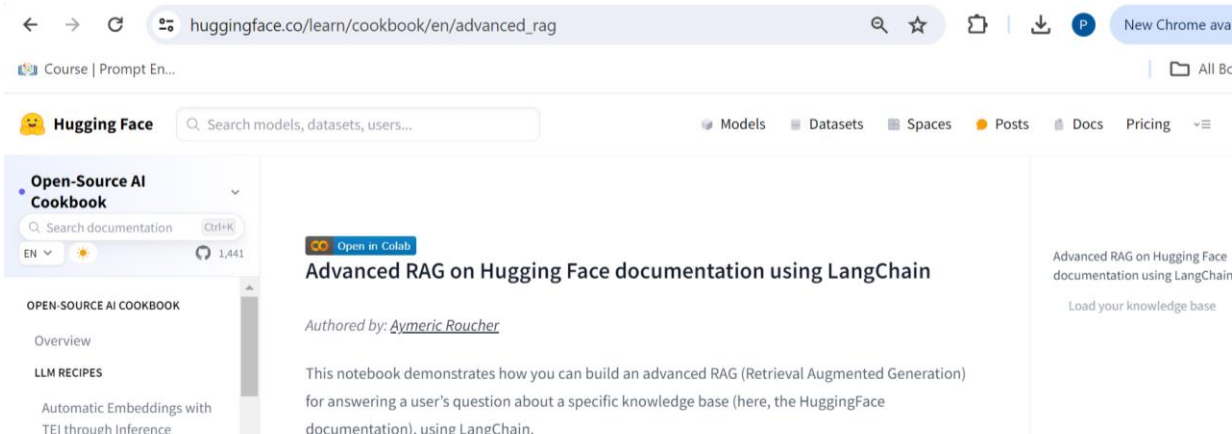
# Case study: Jupyter notebook LLM tutorials (with pip install) on Expanse

- Start with Huggingface/Langchain for RAG

[https://huggingface.co/learn/cookbook/en/advanced\\_rag](https://huggingface.co/learn/cookbook/en/advanced_rag)

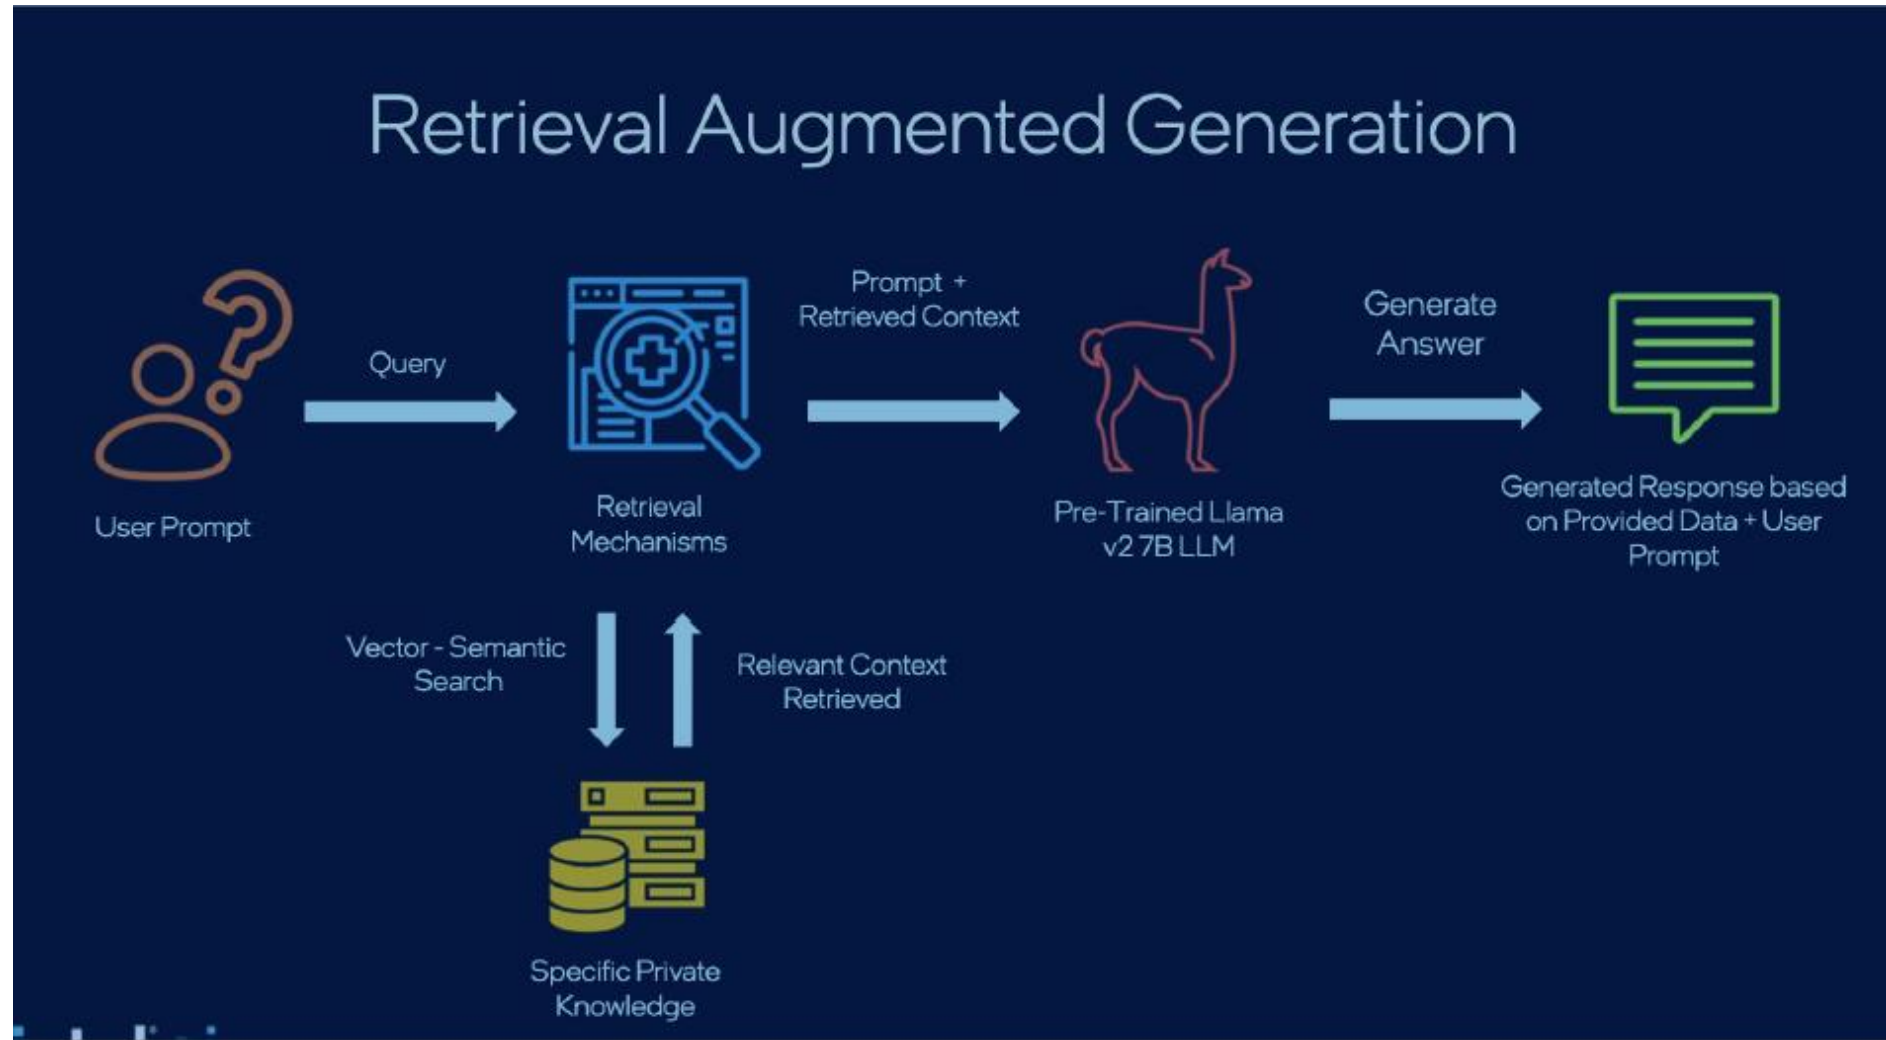
- But use a URL about SLURM for the raw documents

[https://python.langchain.com/v0.2/docs/integrations/document\\_loaders/url/](https://python.langchain.com/v0.2/docs/integrations/document_loaders/url/)





## RAG overview (from the Mai's slide)



# Running notebook tutorials on Expanse

- On Expanse we can run a python singularity container that starts up Jupyter notebooks on a GPU and run !pip install commands  
Pro: easy to follow instructions  
Con: some things not scalable for long workflows

# Running notebook tutorials on Expanse

- On Expanse we can run a python singularity container that starts up Jupyter notebooks on a GPU and run !pip install commands

Pro: easy to follow instructions

Con: some things not scalable for long workflows

OR

- On Expanse we can get a GPU node and use the command line to run pip install commands

Pro: you get to see how/where everything is set up

Con: more typing

added benefit: directly translates to a batch script

# Hugging Face set up using a notebook

- Many tutorials use Jupyter notebooks with “!pip install” commands

```
In [ ]: ▶ !pip install --upgrade 'huggingface_hub[pytorch,cli]'  
        !pip install transformers  
        !pip install accelerate
```

```
In [ ]: ▶ #restart kernel to reload .local modules after installing hugging face hub  
        import huggingface_hub  
        from transformers import AutoTokenizer  
        import transformers  
        import torch
```

*Packages will be in  
/home/userid/.local*

*After installing (~5 min)  
restart kernel to reset  
variables*

# Hugging Face set up using a notebook

- Many tutorials use Jupyter notebooks with “!pip install” commands

```
In [ ]: ▶ !pip install --upgrade 'huggingface_hub[pytorch,cli]'  
!pip install transformers  
!pip install accelerate
```

```
In [ ]: ▶ #restart kernel to reload .local modules after installing hugging face hub  
import huggingface_hub  
from transformers import AutoTokenizer  
import transformers  
import torch
```

*Packages will be in  
/home/userid/.local*

*After installing, restart  
kernel to reset variables*

*BEWARE that “pip install” :*

*might create conflicts with other installs (try clear out/save previous ‘.local’)*

*might try to uninstall a package on the system’s Path path (try the –ignore-installed option)*

*might take a while*

*Otherwise it’s easy*

# Hugging Face set up using the command line (instead of notebook)

## Main Steps

1. Request a GPU node
2. ssh into that node
3. Load modules
4. Run “singularity shell” command
5. Run “pip install” commands
6. Login to hugging face
7. Run your script

*With practice you might prefer to work from command line to see all the details*

# Hugging Face set up using the command line

## Main Steps

### 1. Request a GPU node

```
[train113@login02 ciml-summer-institute-2024]$  
[train113@login02 ciml-summer-institute-2024]$ jupyter-gpu-shared-pytorch
```

### 2. ssh into that node

```
[p4rodrig@login01 HFace]$ squeue -u p4rodrig  
      JOBID PARTITION     NAME     USER ST       TIME  NODES NODELIST(REASON)  
      31537549 gpu-debug  hface-gp p4rodrig  R        0:01      1 exp-7-59  
[p4rodrig@login01 HFace]$  
[p4rodrig@login01 HFace]$  
[p4rodrig@login01 HFace]$ ssh exp-7-59
```

# Hugging Face set up using the command line

## Main Steps

### 3. Load modules

```
[p4rodrig@exp-7-59 HFace]$ module purge
[p4rodrig@exp-7-59 HFace]$ module load gpu
[p4rodrig@exp-7-59 HFace]$ module load slurm
[p4rodrig@exp-7-59 HFace]$
[p4rodrig@exp-7-59 HFace]$ module load singularitypro/3.11
```

### 4. Run “singularity shell” command

```
[p4rodrig@exp-7-59 ~]$ singularity shell /cm/shared/apps/containers/singularity/pytorch
h-latest.sif --nv --bind /expance,/scratch
Singularity>
```



# Hugging Face set up using the command line

## Main Steps

### 5. Run “pip install” commands (takes ~5minutes)

```
Singularity> pip install --upgrade huggingface_hub[pytorch,cli] transformers  
ts  
pip install --upgrade langchain sentence-transformers langchain-community  
pip install --upgrade bitsandbytes pypdf faiss-gpu pydantic
```

### 6. Login to hugging face

```
Singularity> /home/p4rodrig/.local/bin/huggingface-cli login --token hf_cx0F
```

### 7. Run your script

```
Singularity> python3 Rag_example_v1.py > rag_stdout.txt
```

# Hugging Face set up using the command line

## Other Steps

Review output:

```
Now here is the question you need to answer.  
  
Question: How to create a slurm job?<|eot_id|><|start_header_id|>assistant<|end_header_id|>  
  
According to the provided documents, to create a Slurm job, you can use the `sbatch` command. This command is used to submit a job script for later execution. The script will typically contain one or more `srun` commands to launch parallel tasks.  
  
Here is an example of how to use `sbatch`:  
  
`sbatch [options] script.sh`  
  
Where `script.sh` is the name of the job script file, and `[options]` are optional parameters that can be used to specify various settings for the job.  
  
For more information on how to use `sbatch` and other Slurm commands, you can refer to the Slurm documentation, which is available in the provided documents.
```

# Hugging Face set up using the command line

## Other Steps

Review cached models in  
~/.cach/huggingface/hub

```
[p4rodrig@login02 hub]$ du -sh *
417M    models--bert-base-cased
572K    models--bert-base-uncased
512     models----home--p4rodrig--HW2-tests--Yelp1m--bert-it-
-bert-it-vocab.txt
13G     models--meta-llama--Llama-2-7b-chat-hf
15G     models--meta-llama--Meta-Llama-3-8B-Instruct
512     version.txt
[p4rodrig@login02 hub]$
```

# Hugging Face in a batch script

## Other Steps

move .local to new folder, and export PYTHONPATH to avoid potential conflicts

Use “singularity exec” to launch commands in a batch script

```
# Set up paths for python to find packages
export PYTHONPATH=/home/$USER/Local_HFgpu-latest/lib/python3.1
export PATH=/home/$USER/Local_HFgpu-latest/local/bin:$PATH
echo "----- paths -----"
echo $PYTHONPATH
echo $PATH

#You can run hugging face login first (it will put the token i
# and also run it within singularity (b/c it was installed tha
singularity exec --nv --bind /exppanse,/scratch /cm/shared/apps
pytorch-latest.sif /home/$USER/Local_HFgpu-latest/bin/huggingf

#Now you can run the rag example
singularity exec --nv --bind /exppanse,/scratch /cm/shared/apps
pytorch-latest.sif python3 Rag_example_v1.py > rag_stdout.txt
```

# Text for command lines

- The above instructions are in “CommandLineTexts2.txt” file

```
[p4rodrig@login02 HFace]$ more CommandLineTexts.txt
1 Request a GPU node
/cm/shared/apps/sdsc/galyleo/galyleo.sh launch -A use300 -p gpu-debug -n 1 -l
  -G 1 -t 00:30:00 -e singularitypro/3.11 --nv --bind /expanse,/scratch -s /c
ared/apps/containers/singularity/pytorch/pytorch-latest.sif
NOTE: change use300 to quest
```

- Also look at the run-pyt-gpu.. batch script

```
[p4rodrig@login02 HFace]$ more run-pyt-gpudebug-latest.sh
#!/usr/bin/env bash
#SBATCH --job-name=hface-gpu
#SBATCH --account=use300
# -----
#SBATCH --partition=gpu-debug
```

# Hugging Face set up using the command line

Other details like:

The command line command to login to hugging face

```
Singularity> /home/p4rodrig/.local/bin/huggingface-cli login --token hf_cx0F
```

Review cached models in  
~/.cach/huggingface/hub

```
[p4rodrig@login02 hub]$ du -sh *
417M    models--bert-base-cased
572K    models--bert-base-uncased
512     models---home--p4rodrig--HW2-tests--Yelp1m--bert-it-
-bert-it-vocab.txt
13G     models--meta-llama--Llama-2-7b-chat-hf
15G     models--meta-llama--Meta-Llama-3-8B-Instruct
512     version.txt
[p4rodrig@login02 hub]$
```

**end**